

ALEXANDRE TEIXEIRA

AI Engineer · GenAI & Production LLM Systems

alexandretei13@gmail.com ♦ +33 7 86 01 55 04 ♦ [LinkedIn](#) ♦ [GitHub](#) ♦ [Portfolio](#) ♦ Paris, France

0 → 1 AI Engineer — I turn complex business problems into AI products that get deployed and adopted. My approach: understand the workflow, simplify the problem, and build from idea to impact.

EDUCATION

University of Technology of Compiègne (UTC)

Master's-level Engineering Degree (Diplôme d'Ingénieur), Computer Science

2022 – 2026

Compiègne

- Artificial Intelligence, optimization, operations research, logic & search-based problem solving, cybersecurity.

Shanghai University — UTSEUS

Exchange Program in Artificial Intelligence

Sept. 2024 – Jan. 2025

Shanghai

- Natural Language Processing, Generative AI, Data Analysis, Data Visualization.

PROFESSIONAL EXPERIENCE

AI Engineer, SNCF · Internship

Multimodal AI Platform for Meeting Summarization

Sept. 2025 – Feb. 2026

Paris

- Built an internal multimodal platform used across teams for RFIs, video calls and in-room meetings.
- **90% of audio processed in < 5 min**: transcription, speaker diarization and summarization.
- Detected key slides + VLMs to automatically inject visual context into generated summaries.
- Built 0 → prod (React, Python, Azure AI Services, Docker, Jenkins).
- Led an RFI across **15 AI vendors** (Mistral AI, IBM, LinkUp, Vespa, and others); supported the internal rollout of **Mistral Code** for engineering teams.

AI Engineer — Junior-Enterprise Mission, Renault / Alpine

AI Whiteboard for Scaling FMEA Workflows

Mar. 2025 – Jun. 2025

Compiègne

- Built an AI-augmented collaborative whiteboard to generate and structure FMEAs at scale.
- LLM + RAG to propose risks, causes, effects and actions while keeping expert validation in the loop.
- Industrialized inference on an **8× NVIDIA A100** cluster with vLLM and load balancing.

AI Engineer, SNCF · Internship

GenAI Assistant Automating L1 DevOps Support

Sept. 2023 – Feb. 2024

Lyon

- Built the RAG assistant **solo from 0 → 1**; functional MVP shipped in **3 months**, architecture to deployment.
- **82% of L1 ServiceNow tickets resolved automatically**, eliminating most manual first-line handling.
- Product industrialized: **3,000+ monthly users**, 25-person team; featured on BFM Business.

SELECTED PROJECTS

CNN vs. Transformer Architectures for 3D Medical Imaging — with UCL

University College London

Feb. 2026 – Present

- Benchmarking CNN / Transformer models on performance, complexity and inference constraints.

On-Prem LLM Platform for UTC Students

Feb. – Jul. 2025

- Self-hosted LLM platform (OpenWebUI + Ollama, RTX 5090) on UTC's servers, open to all students.

AlphaZero from Scratch — Gopher & Dodo Games

2024

- MCTS, self-play and deep reinforcement learning; PyTorch residual network with policy and value heads.

TECHNICAL SKILLS

AI / LLM Engineering: PyTorch, TensorFlow, Hugging Face, vLLM, LangChain, RAG, embeddings, vector databases, VLMs

Machine Learning: CNNs, Transformers, ResNet, Scikit-learn, model training & evaluation

Cloud / DevOps: Azure AI Services, Azure OpenAI, Docker, Kubernetes, Jenkins, Harbor, Artifactory, GitLab CI/CD

Software: Python, React, FastAPI, Flask, Pydantic, PostgreSQL, REST APIs, SQL, Pandas, NumPy

LANGUAGES & LEADERSHIP

Languages: French (native) · English (professional) · Chinese (basic)

Partnership Manager, Imaginarium Festival (2022–2023) — managed sponsorships for a €600,000-budget student festival.

Founder Breakfast, Shanghai (2024–2025) — engaged with startup founders on AI-driven innovation.