

ALEXANDRE TEIXEIRA

AI Engineer

alexandretei13@gmail.com ♦ +33 7 86 01 55 04 ♦ [LinkedIn](#) ♦ [GitHub](#) ♦ [Portfolio](#) ♦ Paris, France

0→1 AI Engineer — I turn complex business problem statements into AI products that get deployed and adopted. My approach: understand the workflow, simplify the problem, and build from idea to impact.

EDUCATION

- **University of Technology of Compiègne (UTC)** 2022 – 2026
Master's-level Engineering Degree (Diplôme d'Ingénieur), Computer Science Compiègne
 - Artificial Intelligence, optimization, operations research, logic & search-based problem solving.
 - Additional coursework: Psychology of Learning, Project Management, and Responsible Innovation Marketing.
- **Shanghai University — UTSEUS** 2024 – 2025
Exchange Program in Artificial Intelligence Shanghai
 - Natural Language Processing, Generative AI, Data Analysis, Data Visualization.

PROFESSIONAL EXPERIENCE

- **AI Engineer, SNCF · Internship** Sept. 2025 – Feb. 2026
Multimodal AI Platform for Meeting Summarization Paris
 - RFIs, video calls and in-room meetings were buried in hours of audio no team had time to review.
 - Built the **platform 0→prod**: transcription, speaker diarization, summarization, plus key-slide detection and VLMs for visual context (React, Python, Azure AI Services, Docker, Jenkins).
 - 90% of audio processed in < 5 min; adopted across teams.
 - Technical reference for the team on AI direction: **led an RFI** across 15 AI vendors (**Mistral AI, IBM, LinkUp, Vespa**) and the internal rollout of **Mistral Code**.
- **AI Engineer — Junior-Enterprise Mission, Renault / Alpine** Mar. 2025 – Jun. 2025
AI Whiteboard for Scaling FMEA Workflows Compiègne
 - FMEA risk analysis was manual and expert-bound, capping how many workflows could be covered.
 - Built an AI-augmented collaborative whiteboard: LLM + RAG propose risks, causes and actions, experts validate.
 - Industrialized inference on an **8× NVIDIA A100** cluster with vLLM and load balancing.
- **AI Engineer, SNCF · Internship** Sept. 2023 – Feb. 2024
GenAI Assistant Automating L1 DevOps Support Lyon
 - The group was exploring where GenAI could pay off; I found the need in L1 DevOps support, where repetitive ServiceNow tickets were handled manually.
 - Scoped the use case and built the RAG assistant **solo from 0→1**, architecture to deployment; MVP in 3 months.
 - 82% of L1 tickets resolved automatically; **3,000+ monthly users**, 25-person team, featured on BFM Business.
 - The product then grew group-wide **into SNCF GPT**, still running today on the architecture I designed.

SELECTED PROJECTS

- **CNN vs. Transformer Architectures for 3D Medical Imaging — with UCL** Feb. 2026 – Present
University College London — benchmarking performance, complexity and inference constraints.
- **On-Prem LLM Platform for UTC Students** Feb. – Jul. 2025
Self-hosted LLM platform (OpenWebUI + Ollama, RTX 5090) on UTC's servers, open to all students.
- **AlphaZero from Scratch — Gopher & Dodo Games** 2024
MCTS, self-play and deep RL; PyTorch residual network with policy and value heads.

TECHNICAL SKILLS

AI / LLM Engineering: PyTorch, Hugging Face, vLLM, LangChain, RAG, embeddings, vector databases, VLMs
Machine Learning: CNNs, Transformers, ResNet, Scikit-learn, model training & evaluation
Cloud / DevOps: Azure AI Services, Azure OpenAI, Docker, Kubernetes, Jenkins, GitLab CI/CD
Software: Python, React, FastAPI, PostgreSQL, REST APIs, SQL, Pandas, NumPy

LANGUAGES & LEADERSHIP

- **Languages** — French (native) · English (professional) · Chinese (basic).
- **Partnership Manager, Imaginarium Festival** — sponsorships for a 600,000 EUR student festival. 2022 – 2023
- **Founder Breakfast, Shanghai** — engaged with startup founders on AI-driven innovation. 2024 – 2025